

A Fail-Closed Certificate Architecture for Literal Gaussian-Prime-Ideal Ownership in a Level-(3) Bianchi Flow

Liang Wang

School of Artificial Intelligence and Automation, Huazhong University of Science and Technology

Luoyu Road 1037, 430070, Hubei, P.R. China

wangliang.f@gmail.com

2 September 2026

Abstract

This article formulates a fail-closed certificate architecture for asking whether primitive loxodromic owners in a fixed level-(3) Gaussian Bianchi geodesic flow can receive one literal Gaussian prime ideal. It is a review-adjudicated synthesis of a frozen 22-source corpus, rather than a theorem of ownership, a mechanism implementation, a quotient computation, or a finite-refinement experiment. The corpus identifies the Picard/Bianchi object, primitive and inversion-aware terminology, adjacent arithmetic-group algorithms, and exact ideal-arithmetic components. It does not provide either the project-specific owner law or the complete primitive-root, conjugacy, inversion, and canonical-unoriented quotient required for evaluation. The literal one-ideal codomain is therefore treated as a deliberately strict stress test, not as a canonical or literature-forced output. We separate two open interfaces: Gate M for mechanism admissibility and Gate Q for quotient completeness. Their independence means only that neither gate entails the other at the design level. Downstream collision performance is firewalled from selecting or repairing either interface. The resulting prospective protocol specifies typed serializers, positive and negative certificates, immutable mechanism registration, controlled stop states, and independent replay. Because every inherited citation lacks a passage locator, claim-to-passage support remains inconclusive. No owner, quotient, statistic, or route advancement is reported.

Keywords: Bianchi geodesic flow; Gaussian prime ideals; primitive conjugacy classes; certificate architecture; strict codomain; evidence synthesis

Traditional Chinese Abstract / 繁體中文摘要

本文提出一套失敗即封閉的證書架構，用以檢驗固定層級三之高斯 Bianchi 測地流中，原始雙曲所有者是否能被一致地指派一個字面上的高斯素理想。研究僅綜整已凍結的二十二項來源及既有評審紀錄，不執行新定理證明、機制建置、商集計算或有限精化實驗。文獻足以界定 Picard/Bianchi 幾何對象、原始與反向辨識語彙、相鄰算術群演算法及理想算術工具，卻未提供本計畫專屬的所有者法則，也未完成最大根、共軛、取逆及無向正規化所需的完整商集。故單一字面素理想被明確視為刻意嚴格的壓力測試，而非典範、伽羅瓦穩定或由文獻必然導出的值域。本文把機制可容許性與商集完備性分成兩道仍開放的關卡；兩者僅在設計蘊涵上互不替代，下游碰撞分離表現不得用來選擇或修補任一關卡。架構進一步規定型別化序列化、正反證書、不可變註冊、停止狀態與獨立重播。所有引文均缺段落定位，故支持狀態仍為不確定；本文不宣稱任何所有者、商集、統計值或路線晉級。

關鍵詞： Bianchi 測地流；高斯素理想；原始共軛類；證書架構；嚴格值域；證據綜整

1 Introduction, Research Question, and Contribution Boundary

The study fixes one mathematical frame before performance is visible. Its dynamical object is the inherited unit-speed geodesic flow on a torsion-free level-(3) Gaussian Bianchi manifold. Hyperbolic arclength is the clock. A prospective owner is a primitive loxodromic conjugacy class; inversion identifies the primary unoriented owner, and powers are repetitions rather than new owners. The requested arithmetic output is one literal nonzero prime ideal of $\mathbb{Z}[i]$. Associates denote the same ideal, while conjugate prime ideals above a split rational prime remain distinct. These are frozen design definitions, not findings extracted from the corpus.

The research question is deliberately conditional: under that strict output type, what must be certified before a performance-independently registered owner mechanism could be evaluated on a certified finite owner population? The question does not assume that a lawful mechanism exists. It also does not ask whether a norm, rational prime, composite ideal, residue class, or unordered conjugate pair might be a useful alternative object. Those alternatives are relevant to the interpretation of the frame, but substituting one of them after review would answer a different question.

The central distinction is between two pre-performance obligations. The mechanism-admissibility obligation asks whether a frozen formula maps each admissible primitive unoriented owner to one literal Gaussian prime ideal while respecting conjugacy, inversion, and repetition. The quotient-completeness obligation asks whether the inherited finite matrix population has been converted into exact primitive roots and canonical unoriented classes with complete positive and negative certificates. They are design-level independent: a formula cannot decide the quotient merely by producing labels, and a complete quotient cannot invent or validate a formula. This is not a proposition that the two failure events are statistically independent, nor a theorem that every possible construction separates into these modules.

The article's contribution is accordingly procedural and falsifiable. It specifies a fail-closed sequence in which neither classification performance nor collision separation can select a mechanism, repair an invalid codomain map, or excuse an incomplete owner quotient. A mechanism can fail; a quotient can remain unresolved; the downstream estimand can remain not evaluable. These are valid outcomes. The architecture is not claimed to be novel, and it is not presented as an existence or nonexistence theorem about intrinsic arithmetic ownership.

1.1 Claim classes and failure semantics

Four kinds of statements must remain distinct throughout the article. First, the frozen object definitions specify what this study elects to mean by flow, clock, primitive owner, repetition, inversion, and arithmetic output. They make the question determinate, but they do not show that an owner mechanism exists. Second, corpus findings describe only what the bounded evidence package can support at its present resolution. A source may identify the relevant group, discuss primitive geodesics, or provide an adjacent algorithm without instantiating the map required here. Third, architectural requirements state what a future implementation would have to serialize and verify. They are normative proof obligations, not completed constructions. Fourth, status statements record what has actually happened: no mechanism has been admitted, no quotient has been certified, and no downstream statistic is available.

This classification prevents several common but consequential substitutions. A well-defined ideal factorization algorithm does not select an ideal for a geometric owner. A conjugacy routine for a neighboring input class is not automatically a solver for the frozen level-(3) relation. An invariant observed on matrix rows does not turn those rows into primitive unoriented owners. Conversely, failure to find a mechanism in a bounded search would not prove nonexistence. Each statement must be evaluated against its own input type, output type, witness, and quantifier.

The literal single-ideal codomain is kept fixed precisely because it is a stress test. The paper asks

whether an unoriented primitive owner contains enough representative-independent information to select one literal Gaussian prime ideal, with associates identified but conjugate split primes still distinct. A later choice to return a norm, rational prime, composite ideal, residue class, or unordered conjugate pair may define an interesting study, but it would alter the frozen output type. Such a change cannot be treated as a repair after the present mechanism or performance has been inspected. Keeping the strict target does not elevate it to a canonical one; it preserves the falsifiability of this particular question.

Fail-closed semantics give informative meaning to an incomplete pipeline. Missing data do not default to a negative theorem, and an unresolved relation does not default to inequality. Gate M can remain unevaluable because no formula has been frozen, or it can refute a particular frozen formula by a reproducible law violation. Gate Q can remain unevaluable because its exact decision procedures are absent, or stop unresolved under a prospectively fixed resource limit. These outcomes are more precise than a generic failure label. They identify which interface lacks evidence without granting conclusions about other mechanisms, broader codomains, or the global arithmetic of closed geodesics.

Reader map. The prospective dependency chain is typed as follows. An *input row* is a serialized projective matrix with subgroup and loxodromic-eligibility evidence. Gate Q may turn eligible rows into oriented primitive classes only after complete root and conjugacy decisions; quotienting those classes by inversion would then define the unoriented *owners*. Gate M may consume only such owners together with one immutable candidate-mechanism registry and may emit either a validated literal ideal or a typed stop receipt. A *performance ledger* and its estimand may be constructed only from a reconciled Gate-Q population and a Gate-M-admissible registry. A missing or failed node prohibits every downstream use, and every node named here remains prospective.

2 Frozen Literature and Theoretical Frame

2.1 The Picard/Bianchi object does not itself supply an owner law

Within the frozen corpus, prime-geodesic work on the Picard manifold is used only to identify the Gaussian quotient and its analytic setting. No theorem, page, section, or paragraph locator has been passage-adjudicated for the present use; support therefore remains INCONCLUSIVE, and this citation may not be transferred to a class-to-ideal construction or owner law. [3].

Within the frozen corpus, this prime-geodesic study is used only as analytic context for the same Gaussian group and its spectral exponential sums. No exact passage locator has been adjudicated; support remains INCONCLUSIVE, and the source is not used to define the project's owner map. [14].

Within the frozen corpus, this source is used only for Gaussian-arithmetic vocabulary in the Picard setting. No exact passage locator has been adjudicated; support remains INCONCLUSIVE, and shared arithmetic vocabulary is not evidence for a unique literal-ideal assignment. [4].

Within the frozen corpus, this source is used only to motivate primitive closed-geodesic, conjugacy-class, inversion-aware, and unoriented terminology in a nearby arithmetic-hyperbolic setting. No exact passage locator has been adjudicated; support remains INCONCLUSIVE, and no decision procedure is transferred to the frozen noncompact level-(3) relation. [8].

Within the frozen corpus, this earlier Picard prime-geodesic work is used only as historical analytic context. No exact passage locator has been adjudicated; support remains INCONCLUSIVE, and the citation supplies neither an owner law nor a literal-ideal selection mechanism. [15].

Within the frozen corpus, this explicit-formula work is used only to broaden the analytic hyperbolic-three-manifold setting. No exact passage locator has been adjudicated; support remains INCONCLUSIVE, and no project mechanism or owner assignment is inferred. [18].

Within the frozen corpus, this original hyperbolic-three-manifold article is retained only under its correction-and-addendum provenance constraint. No exact affected passage has been adjudicated; support remains INCONCLUSIVE, and the uncorrected record may not be transferred to any affected premise or owner mechanism. [1].

Within the frozen corpus, this paired correction is used only to constrain the admissible use of its companion article. No exact corrected passage has been adjudicated; support remains INCONCLUSIVE, and the correction is not treated as a mechanism source or retraction-clearance record. [2].

Within the frozen corpus, this first-party preprint is used only as current context for the exact Gaussian quotient. No exact passage locator has been adjudicated and the record remains a preprint; support is INCONCLUSIVE, with no transfer to the mechanism or quotient interfaces. [21].

Within the frozen corpus, this arithmetic-hyperbolic background is used only for structural vocabulary concerning Kleinian groups, trace fields, and hyperbolic three-manifolds. No exact passage locator has been adjudicated; support remains INCONCLUSIVE, and no direct owner formula is imported. [17].

Within the frozen corpus, this source is used only for Bianchi/Picard actions and imaginary-quadratic arithmetic as structural setting. No exact passage locator has been adjudicated; support remains INCONCLUSIVE, and the setting is not evidence that a literal ideal owner is canonical. [10].

Taken together, these records support the registered distinction between knowing the object and possessing the project-specific owner law.

2.2 Quotient completeness is a composite certificate target

Within the frozen corpus, arithmetic-Kleinian fundamental-domain computation is used only to motivate exact object representations and reduction witnesses. No exact theorem or passage locator has been adjudicated; support remains INCONCLUSIVE, and no implemented level-(3) primitive-unoriented quotient is transferred. [19].

Within the frozen corpus, these general arithmetic-group algorithms are used only to motivate decidability and certificate questions. No exact theorem or passage locator has been adjudicated; support remains INCONCLUSIVE, and the exact project relation, termination proof, and serialization are not imported. [12].

Within the frozen corpus, this primary conjugacy methodology is used only as an adjacent algorithmic component for certain arithmetic groups. No exact passage locator or reduction to the frozen subgroup has been adjudicated; support remains INCONCLUSIVE, and no instantiated level-(3) solver is claimed. [13].

Within the frozen corpus, these presentations are used only to motivate a possible word-to-matrix interface for special linear and Bianchi cases. No exact passage locator has been adjudicated; support remains INCONCLUSIVE, and maximal-root, conjugacy, inversion, and canonicalization certificates remain separate obligations. [22].

Within the frozen corpus, this oriented-geometrizable three-manifold conjugacy algorithm is used only to illustrate a nearby but differently typed hypothesis class. No exact passage locator or reduction to the frozen group has been adjudicated; support remains INCONCLUSIVE, and the algorithm is not transferred. [20].

Within the frozen corpus, this word-hyperbolic conjugacy result is used only to illustrate a conditional decidability framework. No exact passage locator or proof that the frozen group meets the relevant hypothesis has been adjudicated; support remains INCONCLUSIVE, and no project certificate follows. [11].

Within the frozen corpus, constructive integral-matrix conjugacy is used only to motivate witness design. No exact passage locator or reduction has been adjudicated; support remains INCONCLUSIVE, and

ambient equality or conjugacy in $GL(n, \mathbb{Z})$ is not transferred to the level-(3) projective subgroup relation. [9].

The quotient target therefore remains a composition problem. It requires an exact projective-lift convention, loxodromic eligibility, maximal group-element roots, level-(3) conjugacy, inversion links, deterministic canonical unoriented identifiers, transitivity checks, and population accounting. A bounded search that does not find a conjugator is not a negative certificate. The corpus supplies possible ingredients but no one frozen source supplies this complete target.

2.3 Ideal arithmetic validates outputs but does not choose owners

Within the frozen corpus, these number-field and ideal-factorization algorithms are used only to motivate validation of nonzero Gaussian prime ideals and rational-prime splitting. No exact passage locator has been adjudicated; support remains INCONCLUSIVE, and the algorithms do not select an ideal owner. [7].

Within the frozen corpus, effective algebraic-number-theory methods are used only to motivate exact arithmetic and certificate-oriented implementation. No exact passage locator has been adjudicated; support remains INCONCLUSIVE, and no class-to-ideal selection law is imported. [16].

Within the frozen corpus, practical computational-number-theory interfaces are used only to motivate serialization and validation requirements. No exact passage locator has been adjudicated; support remains INCONCLUSIVE, and the interfaces do not justify assigning one branch of a split prime to an unoriented owner. [5].

Within the frozen corpus, maximal-power detection for integral ideals is used only to motivate an ideal-side witness type. No exact passage locator has been adjudicated; support remains INCONCLUSIVE, and ideal powers are not transferred to maximal roots of group elements or geodesic repetition. [6].

The literal codomain is a deliberately strict frozen frame, not a canonical or literature-forced target. For a registered candidate M on eligible representatives, Gate M requires the exact laws $M(h\gamma h^{-1}) = M(\gamma)$ for the frozen subgroup relation and $M(\gamma^{-1}) = M(\gamma)$ for descent to an unoriented owner. Gaussian conjugation $\mathfrak{p} \mapsto \bar{\mathfrak{p}}$ is a separate codomain involution. Thus a candidate that instead yields $M(\gamma^{-1}) = \overline{M(\gamma)}$ fails unoriented descent when the two split prime ideals are distinct. Such a failure refutes only that registered candidate under the literal single-ideal codomain; it does not rule out other mechanisms or broader codomains.

3 Executed Methodology: Frozen Evidence Synthesis Only

The executed evidence synthesis was bounded to 48 record manifestations. Deduplication removed 12 duplicate manifestations; 36 unique records were screened at title or abstract level, 22 were admitted after scope screening, and 17 admitted records entered the workflow's peer-reviewed journal-or-correction numerator. The frozen inventory names P29-S01 through P29-S22, and independent verification records 22 VERIFIED existence outcomes with bounded metadata normalization. The accompanying evidence files preserve the queried interfaces, run dates, exact query strings, normalization rules, deduplication rule, admitted identifiers, and source-role coding. They do not preserve row-level identifiers and decisions for every screened-out manifestation, so the aggregate counts are auditable but a complete external-search replay is unavailable.

Each admitted record was encoded in a source-effect matrix with support class, claim-fitness grade, admissible use, prohibited stronger use, compatibility role, and locator limitation. The matrix was synthesized into the research report and then independently assessed from editorial, domain, methodology, and adversarial perspectives before author adjudication. This reader-facing account concerns evidence selection and synthesis only. It retrieved or created no new database record, publisher page, full text, passage, quotation, scientific artifact, experiment, or result.

The method has an explicit evidentiary ceiling. In this paper, all 22 inherited prose citation pairs retain `anchor : none`; across the five-paper Round-10 batch, all 144 inherited citation pairs likewise have no passage locator. Exact theorem pages, sections, paragraph identifiers, and quotation spans were not frozen. Source identity, DOI/reference closure, metadata verification, and bounded abstract or scope inspection do not establish claim-to-passage faithfulness. The claim-to-passage status for every cited premise in this article is therefore `INCONCLUSIVE`. No direct quotation appears.

Correction, source-status, and integrity checks remain separately typed. P29-S06 stays bound to P29-S07 for affected use. P29-S09 remains a preprint. Structured retraction status is `NOT_CHECKED` for all 22 source rows, and source-level conflict-of-interest status is `UNKNOWN_NOT_AUDITED`. The absence of a fabricated or unverifiable inventory row is not retraction clearance, and a correction binding is not a full theorem-applicability audit.

3.1 Transformation from evidence records to manuscript claims

Stage 2 converts the frozen evidence synthesis into a publication-shaped argument without increasing its scientific strength. Each registered claim has a preassigned source set and an explicit negative constraint. The manuscript may narrow a claim, explain its logic, or state a prospective certificate field, but it may not add an unassigned source, infer a stronger theorem, or turn an intended interface into an executed result. This rule is especially important where several adjacent sources appear to compose naturally: composition at the level of a research plan is not composition at the level of a proved terminating algorithm.

Source identity is also kept separate from passage support. The source inventory, DOI closure, bounded metadata checks, and support-role coding justify citing a record for a narrow contextual purpose. They do not identify an exact theorem statement or discharge its hypotheses. For that reason every citation is paired with the same unresolved passage status. The correction relation between P29-S06 and P29-S07 is retained, and the preprint status of P29-S09 remains visible, but neither condition is converted into a general integrity clearance. This is a deliberate limit of the present manuscript rather than a silent assumption.

The executed outcome of this transformation is therefore modest but concrete: a coherent argument, a closed bibliography with all 22 frozen identifiers, and an auditable set of claims whose prohibited upgrades remain legible. No numerical or algebraic computation is hidden inside the prose conversion. No candidate formula, matrix population, quotient ledger, collision table, or replay receipt was generated. The article can guide a later scientific stage, but it cannot serve as evidence that the later stage has already occurred.

4 Review-Adjudicated Proof and Method Architecture

4.1 Gate M: mechanism admissibility under the strict frame

Gate M is a prospective logical and serialization contract. Before any collision outcome is visible, one candidate mechanism must be frozen by immutable formula bytes, source trace, admissible domain, literal-ideal codomain, normalization, associate convention, and the equations $M(h\gamma h^{-1}) = M(\gamma)$ and $M(\gamma^{-1}) = M(\gamma)$ on the frozen owner relation. The validator must distinguish a nonzero Gaussian prime ideal from an element, generator, norm, rational prime, composite ideal, factorization list, residue, or unordered conjugate pair. It must also record whether Gaussian conjugation fixes or exchanges the output. A split branch must follow from the registered formula rather than from observed collisions, and failure is scoped to that exact formula.

Gate M has an explicit precedence rule. If the candidate violates literal split-branch compatibility under inversion, the primary disposition is `SPLIT_IDEAL_CODOMAIN_OBSTRUCTION` and the same receipt

records `formal_map_refuted=true`. If that predicate passes but another declared formal law fails, the disposition is `FORMAL_MAP_REFUTED`. Only a candidate passing every declared law and output-type check may receive `MECHANISM_ADMISSIBLE`. These outcomes concern the exact registered candidate under the frozen frame, never all possible mechanisms or codomains, and none has been produced here.

4.2 Gate Q: primitive-unoriented quotient completeness

Gate Q is a prospective population-construction contract. Every eligible inherited matrix row would need an exact loxodromic decision, a maximal group-element root witness, a level-(3) conjugacy class assignment, an inversion link, and one canonical unoriented owner ID. Positive relations require reproducible witnesses. Negative relations require a complete decision procedure or an equivalently auditable certificate; exhaustion of a bounded search is not enough. Deterministic serializers must make repeated runs byte-comparable.

The gate remains unresolved independently of Gate M. If exact completeness cannot be established from the object and algorithms, the outcome is `QUOTIENT_NOT_EVALUABLE`. If a predeclared resource limit is reached while relations remain unresolved, the outcome is `QUOTIENT_UNRESOLVED_STOP`. A complete quotient is not inferred from the number of rows, apparent matrix uniqueness, or a mechanism's labels. No quotient computation or certificate exists in this phase.

4.3 Prospective certificate graph and replay obligations

The first prospective node is `ObjectLedger/v1`. Each row must contain canonical projective-matrix bytes, lift and central-sign conventions, subgroup-membership evidence, an exact loxodromic-eligibility disposition, serializer version, and hashes of every dependency. Its validator accepts only schema-conforming rows with reproducible membership and eligibility witnesses; malformed input yields `OBJECT_ROW_MALFORMED`, and an ineligible row yields an explicit exclusion receipt. Positive, negative, malformed, and serialization-round-trip fixtures must carry predeclared expected dispositions. No such ledger or fixture has been executed.

The second prospective node is `QuotientLedger/v1`. It must bind each admitted object hash to maximal-root, subgroup-conjugacy, inversion, transitive-closure, and canonical-owner decisions. Positive relations require replayable witnesses; negative decisions require a complete terminating procedure for the exact frozen relation. The validator rejects missing root exponents, unverifiable conjugators, inconsistent inverse links, nontransitive class IDs, and population totals that do not reconcile inputs, exclusions, repetitions, oriented classes, and unoriented owners. Root, inverse, conjugate, nonconjugate, malformed-witness, and resource-limit fixtures must have frozen expected stop states. No quotient ledger or fixture has been executed.

The third prospective node is `MechanismRegistry/v1`. A record must bind immutable formula bytes, domain, codomain, source trace, normalization, associate convention, split-conjugate treatment, and the exact conjugacy, inversion, and repetition predicates. The validator separately checks literal prime-ideal type, representative independence, and every declared law, and it attaches any failure only to the record hash. Fixtures must include admissible inert outputs, split conjugate pairs, a deliberate inversion-law failure, a nonprime output, a malformed formula record, and a timeout or unavailable-dependency case, each with a predeclared typed disposition. No mechanism record or fixture has been executed.

The fourth prospective node is `PerformanceLedger/v1`. It may consume only the exact `QuotientLedger/v1` population hash and a `MECHANISM_ADMISSIBLE` registry hash. Its schema must freeze baseline collisions, holdout membership, controls, clock, and the definition of $S_H(M)$ before output access. The validator rejects raw-row scores, incomplete quotients, mismatched mechanism bytes, post-result control changes, or unreconciled populations. Fixtures must include a valid empty-result skeleton, a population-hash

mismatch, a mechanism-hash mismatch, and a post-freeze mutation, all without supplying scientific values. No performance ledger or fixture has been executed.

The prospective `IndependentReplayReceipt/v1` must enumerate every consumed content hash, validator version, fixture identifier, expected disposition, observed disposition, and first mismatch. The verifier may reuse parsers and public schemas but must not import producer-side canonicalization, relation-decision, formula-evaluation, or scoring routines; those transformations require independently implemented checks or direct witness verification. Replay must cover positive, negative, malformed, inversion, root, and resource-limit fixtures and must fail closed on an unrecognized proof type. Agreement would certify only the registered finite procedure. No verifier, fixture execution, or receipt exists in the present work.

4.4 The limited meaning of independence

The two gates are independent in an architectural entailment sense only. Closure of Gate M does not entail closure of Gate Q because producing lawful labels does not decide primitive roots, conjugacy, or inversion classes. Closure of Gate Q does not entail closure of Gate M because a certified owner population does not define a literal ideal-selection law. Failure of either gate blocks downstream performance, but their error rates, computational costs, and empirical behavior have not been observed. The present article therefore makes no claim of mathematical, statistical, or empirical independence.

4.5 Performance firewall

Only after both gates pass could a prospectively frozen owner partition expose a finite collision-separation estimand such as $S_H(M)$. The statistic would count only predeclared baseline-collision owner pairs in a certified holdout for which a fixed mechanism returns different valid literal ideals. It could not choose the mechanism, repair a failed law, or compensate for duplicate/repeated owners. No value of S_H is reported. Even a positive future value would be a bounded finite-refinement result rather than a global orbit-to-prime-ideal theorem.

Typed controls would remain separate from the primary score. Permuting ideal labels without changing equivalence classes is a relabeling invariance test, not arithmetic specificity. Replacing literal primes with composite objects is a type violation, not a matched control. Directionally sensitive controls would need to be registered before output access and assessed under the same object, owner, and clock definitions.

Prospective control interpretation. An owner-label permutation preserves the quotient and tests scoring sensitivity to label identity; failure diagnoses label dependence but cannot establish arithmetic specificity. An inversion-paired control preserves the underlying owner population while testing representative independence; failure diagnoses an owner-law or serialization defect but cannot identify the correct ideal branch. A broadened-codomain control, such as an unordered conjugate pair, changes the output type and may diagnose whether a literal-branch obstruction is codomain-specific; it cannot rescue the registered single-ideal mechanism or count as performance under the frozen estimand. Every construction, statistic, and tolerance remains to be registered before output access, and no control has been run.

5 Evidence-Synthesis Findings

The first registered finding is a support boundary. The frozen corpus identifies the Picard/Bianchi geodesic object and supplies primitive, inversion, group-algorithm, and ideal-arithmetic vocabulary. It does not define the project-specific Gaussian-prime-ideal owner law. Object identity and terminology therefore cannot be promoted into an owner map, a quotient, or a refinement score.

The second finding is that quotient completeness remains a composite certificate target. The adjacent methods do not arrive as one theorem or implemented solver for the exact level-(3) relation. The required hypothesis bridges, termination logic, positive and negative witnesses, and deterministic canonicalization remain prospective. This gap is genuine within the frozen design even if the literal codomain were later replaced in a different study.

The third finding is the review-adjudicated codomain interpretation. One literal Gaussian prime ideal is a deliberately strict project frame, not an intrinsic consequence of the geometry or literature. A split-branch obstruction under that frame would be conditional and construct-specific. An unordered Galois orbit or another equivariant output might avoid that problem, but considering it would require a separately frozen research question; it cannot be introduced as a repair after performance access.

The fourth finding is the pre-performance separation of mechanism admissibility and quotient completeness. Their design-level non-entailment prevents performance from selecting or repairing either interface. The contribution is the resulting fail-closed certificate architecture, not a theorem, novelty determination, implemented mechanism, or scientific result.

6 Reproducibility and Prospective Implementation Interface

Reader-auditable evidence is limited to four files at the commit-pinned repository locator https://github.com/maris205/hilbert-polya-structure/tree/5ca65027918c0fce7ef9af82f3faf2e46ed6530c/flow_systems/papers/29-bianchi-ideal-owner-refinement. Their SHA-256 digests are: `notes/stage1_phase2_annotated_bibliography.md`, `c4d71637e5676337326d2eb78dcdd64d78b4b116a397c50c54a081d7c5e2650b`; `notes/stage1_phase2_source_inventory.tsv`, `67ed7713bd6881d11466dc16755c7660a458c52e07ee072d086d6467f8ad7bd`; `notes/stage1_phase2_source_verification.tsv`, `bcf5fa7af07f353fbcaaa6fca319e79f173d7a6af070b582`; and `notes/stage1_phase3_literature_matrix.tsv`, `219309f0c8cea9106ba162c95f2de266b78b1c5d8e0d541`. These files support record identity, bounded screening counts, metadata status, and source-role coding only. Other internal workflow files are provenance surfaces, not an independently replayable scientific package.

A later scientific implementation would expose five closed interfaces: `ObjectLedger/v1`, `QuotientLedger/v1`, `MechanismRegistry/v1`, `PerformanceLedger/v1`, and `IndependentReplayReceipt/v1`. Each interface must publish canonical bytes, dependency hashes, an allowed proof-type registry, validator acceptance predicates, and typed stop states. Producer and verifier may share only schemas, parsers, and witness formats; the verifier may not reuse producer-side canonicalization, relation-decision, mechanism-evaluation, or scoring logic. The required positive, negative, malformed, inversion, root, and resource-limit fixtures are specifications with expected dispositions, not executed tests. Accordingly, no replay, robustness, sensitivity, or usefulness result is reported.

All five interfaces are prospective. No serializer, engine, registry entry, control result, or replay receipt is claimed to exist at Stage 2. Their purpose is to define what evidence would be needed for later evaluation and what stop state must be returned when evidence is missing.

7 Discussion and Implications

The strict-frame clarification changes the interpretation of the first gate without weakening the freeze. The literal codomain can be scientifically useful as a stress test of whether an unoriented geometric owner carries enough intrinsic data to select a non-Galois-stable branch. Yet failure would primarily speak to that deliberately demanding type choice. It must not be advertised as an obstruction to every Galois-stable or orientation-enriched ownership concept.

The quotient gate has a different status. It concerns the validity of the observational unit. Rows cannot become primitive owners through labeling, and finite performance on uncertified rows would not estimate the registered construct. This obligation persists under alternative codomains and is therefore independent of the split-frame question in the limited architectural sense defined above.

The combined architecture encourages interpretable negative outcomes. It distinguishes absence of a frozen admissible mechanism from formal refutation of one mechanism, and it distinguishes an unavailable exact quotient from a resource-bounded unresolved stop. Those distinctions prevent an engineering failure from being overread as a theorem and prevent a finite numerical success from concealing a type error.

The bounded contribution is a project-specific synthesis of typed owner semantics, a literal-ideal stress test, fail-closed Gate-Q and Gate-M contracts, and an independent-replay specification. Certificate methods, proof-carrying computation, and replay-oriented workflow design are therefore comparison classes rather than components claimed as new in isolation. The frozen source study did not perform a field-general novelty or priority review, so no global originality claim is made. Because the prospective schemas, fixtures, controls, and performance ledger have not been executed, practical usefulness and scientific performance also remain unevaluated.

8 Acknowledged Limitations

The source study is bounded rather than exhaustive. The frozen files disclose queried interfaces, run dates, exact query strings, normalization and deduplication rules, admitted identifiers, and source-role decisions. They do not contain row-level identifiers and inclusion or exclusion decisions for every screened-out manifestation or a complete external-interface run log. The admitted 22-source corpus and aggregate counts can therefore be inspected, but the original search cannot be independently replayed record by record. This limitation precludes literature-absence, exhaustive-coverage, and novelty claims.

All 22 paper-specific citation pairs, and all 144 citation pairs in the five-paper batch, lack passage locators. Consequently, statements near theorem boundaries have been narrowed to what the frozen source-verification summaries permit, but claim-to-passage support remains INCONCLUSIVE. No theorem-hypothesis transfer, maximal-root formula, conjugacy procedure, or split-ideal fact is certified at page or theorem level by this article. P29-S06/P29-S07 correction handling remains mandatory, structured retraction screening was not run, and source-level conflicts were not audited.

No separately authorized field-wide novelty analysis or passage adjudication was performed. The contribution is therefore limited to the manuscript’s project-specific synthesis and prospective specification; it is not a priority claim against all certificate-methods, proof-carrying computation, or arithmetic-dynamics workflows. Source-role premises remain INCONCLUSIVE at passage level wherever no exact locator has been supplied, and the machine-audit comments preserve that boundary.

No mechanism, quotient, owner population, collision set, S_H value, control result, or replay receipt exists here. No nonexistence theorem follows from the missing mechanism, and no universal split-prime obstruction is asserted. The literal codomain remains fixed for this study even though it is not claimed to be intrinsically privileged. The exact quotient remains independently unresolved.

The next authorization should also preserve a strict order of access. Passage adjudication should occur before a theorem is mapped to the object; the theorem-to-object map should be frozen before an implementation is evaluated; and both Gate M and Gate Q artifacts should be sealed before collision labels or holdout performance are exposed. This order is not a claim that one implementation strategy is uniquely correct. It is a safeguard against choosing definitions, stopping rules, or split branches because they improve a visible finite score.

A later manuscript could report partial progress without weakening this order. For example, it could close source locators while leaving both scientific gates open, instantiate a complete quotient while

reporting no admissible mechanism, or refute one registered mechanism while leaving the broader search unresolved. Each outcome would require its own hashes and receipts and would retain the strict distinction between failure of a candidate and nonexistence in the whole hypothesis class. Such modular reporting is preferable to a single binary success label because it keeps the mathematical scope of every negative conclusion visible.

9 Future Work

The smallest source-level step would be a separately authorized locator pass for the few premises closest to an owner formula, group-root decision, and literal split-branch law. Any exact passage would need its hypotheses and prohibited transfers frozen; identity verification alone is insufficient.

The smallest object-level step would be a theorem-to-object map for projective lifts, subgroup membership, loxodromic status, maximal roots, level-(3) conjugacy, inversion, and deterministic serialization. It should specify stop states before implementation. A mechanism search and a quotient implementation should remain separate until both are independently frozen.

Only after those steps should a scholar-confirmed scientific design register a mechanism, owner partition, holdout, controls, resource limits, and replay protocol. Alternative Galois-stable or oriented codomains may be studied later as distinct hypotheses, not as post hoc replacements for the frozen literal target.

10 Conclusion

The frozen literature is sufficient to formulate a strict, auditable ownership study but not to produce an owner law or a certified owner population. The present article makes the construct boundary explicit: the literal Gaussian-prime-ideal codomain is deliberately strict and frozen, not canonical or forced, so any split obstruction is conditional on that frame. Quotient completeness is a separate unresolved problem. The two gates are independent only as non-entailing design obligations, and performance is downstream of both.

The article’s defensible result is an evidence-synthesis and certificate architecture with explicit failure states and prospective interfaces. It is not a theorem, scientific computation, novelty finding, or Route advance. The formal Route-A tuple remains UNASSIGNED, positive arithmetic A2 remains 0/1, and Route B remains uninvoked.

The immediate value of the architecture is diagnostic: it prevents a missing owner law, an incomplete observational quotient, and weak finite separation from being collapsed into one ambiguous outcome. A later study can close, refute, or stop at each interface with evidence appropriate to that interface. Until such artifacts exist, Gate M and Gate Q remain open and $S_H(M)$ remains undefined rather than zero.

Author Contributions

Liang Wang is the responsible human author. CRediT-style roles: Conceptualization; Methodology; Writing—review and editing; Project administration. He supplied the conceptual direction, made the author-adjudication decisions, approved the recorded workflow gates, interpreted the bounded evidence synthesis, and accepts final accountability for the manuscript. AI assistance is disclosed separately and is not credited with authorship.

Funding

This research received no funding.

Competing Interests

The author declares no competing interests.

Ethics Statement

Ethics approval, informed consent, and human-subject review are not applicable. This article concerns mathematical objects, published-source records, and project-owned workflow artifacts; it reports no human participants, intervention, identifiable personal data, or human-related dataset.

Data and Materials Availability

The four manuscript-audit files named with SHA-256 digests in the reproducibility section are available at the commit-pinned repository path given there. They support the bounded evidence-synthesis claims stated in this article. Row-level excluded-record decisions, passage-level adjudications, and a complete external-search replay are unavailable. No owner mechanism, quotient ledger, certificate, fixture execution, experiment, or numerical result was created for this revision.

AI-Assistance Disclosure and Verification Limitation

OpenAI Codex, operating in the GPT-5 model family, assisted during the session dated 2 September 2026 UTC; the exact backend snapshot or build was not exposed. Assistance covered organization of the frozen bibliography and source identities, bounded evidence synthesis, drafting, review integration, LaTeX conversion, and deterministic checks of identifiers, citations, references, hashes, and word counts. No AI-authored scientific data, experiment, mechanism, quotient, proof, or numerical result is presented. AI assistance is not credited with authorship.

Liang Wang made the recorded stage-gate confirmations and author-adjudicated framing decisions. Neither that responsibility nor the present composition implies human full-text or source-passage verification. Verification remains limited to frozen metadata, abstracts or authorized scope notes, source-verification ledgers, the source-effect matrix, and review artifacts. Exact passages were not frozen; all citation anchors therefore remain none, and every claim-to-passage assessment remains INCONCLUSIVE pending Stage 2.5.

References

- [1] Muharem Avdispahić. On the prime geodesic theorem for hyperbolic 3-manifolds. *Mathematische Nachrichten*, 291(14–15):2160–2167, 2018. doi: 10.1002/mana.201700190. URL <https://doi.org/10.1002/mana.201700190>.
- [2] Muharem Avdispahić. Errata and addendum to On the prime geodesic theorem for hyperbolic 3-manifolds. *Mathematische Nachrichten*, 292(4):691–693, 2019. doi: 10.1002/mana.201800467. URL <https://doi.org/10.1002/mana.201800467>.

- [3] Olga Balkanova and Dmitry Frolenkov. Prime geodesic theorem for the Picard manifold. *Advances in Mathematics*, 375:107377, 2020. doi: 10.1016/j.aim.2020.107377. URL <https://doi.org/10.1016/j.aim.2020.107377>.
- [4] Antal Balog, András Biró, Giacomo Cherubini, and Niko Laaksonen. Bykovskii-type theorem for the Picard manifold. *International Mathematics Research Notices*, 2022(3):1893–1921, 2022. doi: 10.1093/imrn/rnaa128. URL <https://doi.org/10.1093/imrn/rnaa128>.
- [5] Karim Belabas. Topics in computational algebraic number theory. *Journal de Théorie des Nombres de Bordeaux*, 16(1):19–63, 2004. doi: 10.5802/jtnb.433. URL <https://doi.org/10.5802/jtnb.433>.
- [6] Karim Belabas and Denis Simon. Power detection over number fields. *Mathematics of Computation*, 93(348):1953–1961, 2024. doi: 10.1090/mcom/3913. URL <https://doi.org/10.1090/mcom/3913>.
- [7] Henri Cohen. *A Course in Computational Algebraic Number Theory*, volume 138 of *Graduate Texts in Mathematics*. Springer, 1993. doi: 10.1007/978-3-662-02945-9. URL <https://doi.org/10.1007/978-3-662-02945-9>.
- [8] Lindsay Dever and Djordje Milićević. Ambient prime geodesic theorems on hyperbolic 3-manifolds. *International Mathematics Research Notices*, 2023(1):588–635, 2023. doi: 10.1093/imrn/rnab048. URL <https://doi.org/10.1093/imrn/rnab048>.
- [9] Bettina Eick, Tommy Hofmann, and E. A. O’Brien. The conjugacy problem in $GL(n, \mathbb{Z})$. *Journal of the London Mathematical Society*, 100(3):731–756, 2019. doi: 10.1112/jlms.12246. URL <https://doi.org/10.1112/jlms.12246>.
- [10] Jürgen Elstrodt, Fritz Grunewald, and Jens Mennicke. *Groups Acting on Hyperbolic Space: Harmonic Analysis and Number Theory*. Springer Monographs in Mathematics. Springer, 1998. doi: 10.1007/978-3-662-03626-6. URL <https://doi.org/10.1007/978-3-662-03626-6>.
- [11] David B. A. Epstein and Derek F. Holt. The linearity of the conjugacy problem in word-hyperbolic groups. *International Journal of Algebra and Computation*, 16(2):287–305, 2006. doi: 10.1142/S0218196706002986. URL <https://doi.org/10.1142/S0218196706002986>.
- [12] Fritz Grunewald and Daniel Segal. Some general algorithms. I: Arithmetic groups. *Annals of Mathematics*, 112(3):531–583, 1980. doi: 10.2307/1971091. URL <https://doi.org/10.2307/1971091>.
- [13] Fritz J. Grunewald. Solution of the conjugacy problem in certain arithmetic groups. In S. I. Adian, W. W. Boone, and Graham Higman, editors, *Word Problems II*, volume 95 of *Studies in Logic and the Foundations of Mathematics*, pages 101–139. North-Holland, 1980. doi: 10.1016/S0049-237X(08)71335-1. URL [https://doi.org/10.1016/S0049-237X\(08\)71335-1](https://doi.org/10.1016/S0049-237X(08)71335-1).
- [14] Ikuya Kaneko. The prime geodesic theorem for $PSL_2(\mathbb{Z}[i])$ and spectral exponential sums. *Algebra & Number Theory*, 16(8):1845–1887, 2022. doi: 10.2140/ant.2022.16.1845. URL <https://doi.org/10.2140/ant.2022.16.1845>.
- [15] Shin-ya Koyama. Prime geodesic theorem for the Picard manifold under the mean-Lindelöf hypothesis. *Forum Mathematicum*, 13(6):781–793, 2001. doi: 10.1515/form.2001.034. URL <https://doi.org/10.1515/form.2001.034>.

- [16] H. W. Lenstra, Jr. Algorithms in algebraic number theory. *Bulletin of the American Mathematical Society*, 26(2):211–244, 1992. doi: 10.1090/S0273-0979-1992-00284-7. URL <https://doi.org/10.1090/S0273-0979-1992-00284-7>.
- [17] Colin Maclachlan and Alan W. Reid. *The Arithmetic of Hyperbolic 3-Manifolds*, volume 219 of *Graduate Texts in Mathematics*. Springer, 2003. doi: 10.1007/978-1-4757-6720-9. URL <https://doi.org/10.1007/978-1-4757-6720-9>.
- [18] Maki Nakasuji. Prime geodesic theorem via the explicit formula of Ψ for hyperbolic 3-manifolds. *Proceedings of the Japan Academy, Series A, Mathematical Sciences*, 77(7):130–133, 2001. doi: 10.3792/pjaa.77.130. URL <https://doi.org/10.3792/pjaa.77.130>.
- [19] Aurel Page. Computing arithmetic Kleinian groups. *Mathematics of Computation*, 84(295):2361–2390, 2015. doi: 10.1090/S0025-5718-2015-02939-1. URL <https://doi.org/10.1090/S0025-5718-2015-02939-1>.
- [20] Jean-Philippe Préaux. Conjugacy problem in groups of oriented geometrizable 3-manifolds. *Topology*, 45(1):171–208, 2006. doi: 10.1016/j.top.2005.06.002. URL <https://doi.org/10.1016/j.top.2005.06.002>.
- [21] Zhi Qi. On the symmetric square large sieve for $\mathrm{PSL}_2(\mathbb{Z}[i]) \backslash \mathrm{PSL}_2(\mathbb{C})$ and prime geodesic theorem for $\mathrm{PSL}_2(\mathbb{Z}[i]) \backslash \mathbb{H}^3$, 2024. URL <https://arxiv.org/abs/2407.17959>. Preprint.
- [22] Richard G. Swan. Generators and relations for certain special linear groups. *Advances in Mathematics*, 6(1):1–77, 1971. doi: 10.1016/0001-8708(71)90027-2. URL [https://doi.org/10.1016/0001-8708\(71\)90027-2](https://doi.org/10.1016/0001-8708(71)90027-2).